


Zihao He

[Email](#)[Homepage](#)[Google Scholar](#)

RESEARCH INTERESTS

RLHF, Reward Modeling, LLM Alignment & Safety, LLM Evaluation, Vision Language Models

EXPERIENCE

Meta

Menlo Park, CA

Research Scientist, Monetization GenAI

02/2025 – Present

- Train reward models on large-scale pairwise human preference data for image/video generation; design and run RLHF pipelines to steer generative model behavior via policy gradient methods.
- Build LLM-as-a-Judge evaluation systems for multimodal generation with fine-grained quality rubrics aligned to human ratings.
- Lead data curation pipelines (prompt rewriting, synthetic augmentation, bias detection) and scalable distributed training workflows for vision-language reward and policy models.

TikTok Inc.

San Jose, CA

Research Scientist Intern

05/2023 – 08/2023

- Developed alignment-based music captioning via contrastive audio-text pretraining (published at EMNLP 2023).

Amazon.com Services, Inc.

Seattle, WA (Remote)

Applied Scientist Intern

05/2021 – 08/2021

- Built text-embedding-based book personalization for Amazon Kids.

DiDi Global

Beijing, China

Algorithm Development Intern

04/2019 – 07/2019

- Multi-task learning for multi-label driver misbehavior classification.

EDUCATION

University of Southern California

Los Angeles, CA

Ph.D. in Computer Science

08/2019 – 02/2025

- Advisor: Prof. Kristina Lerman Thesis: *Aligning Large Language Models with Human Perspectives*

Tsinghua University

Beijing, China

M.Eng. in Computer Technologies (Transferred)

09/2018 – 06/2019

Beijing University of Posts and Telecommunications

Beijing, China

B.Eng. in Communications Engineering

09/2014 – 06/2018

SELECTED PUBLICATIONS

- J. Wu, X. Kan, **Z. He**, S. Tan, B. Pan, K. Zhang. “Multi-Task Reinforcement Learning for Enhanced Multimodal LLM-as-a-Judge.” *ACL Industry Track*, 2026. [\[Link\]](#)
- **Z. He**, S. Guo, A. Rao, K. Lerman. “Whose Emotions and Moral Sentiments Do Language Models Reflect?” *Findings of ACL*, 2024. [\[Link\]](#)
- **Z. He**, M.D. Chu, R. Dorn, S. Guo, K. Lerman. “Community-Cross-Instruct: Unsupervised Instruction Generation for Aligning Large Language Models to Online Communities.” *EMNLP*, 2024. [\[Link\]](#)
- K. Chen, **Z. He**, J. Yan, T. Shi, K. Lerman. “How Susceptible are Large Language Models to Ideological Manipulation?” *EMNLP*, 2024. [\[Link\]](#) **Best Paper Runner-up, SeT LLM @ ICLR 2024**

- M.D. Chu, **Z. He**, R. Dorn, K. Lerman. “Improving and Assessing the Fidelity of LLM Alignment to Online Communities.” *NAACL*, 2025. [\[Link\]](#)
- K. Chen, T. Shi, **Z. He**, K. Lerman. “Steer-Bench: A Benchmark for Evaluating the Steerability of Large Language Models.” *EMNLP*, 2025. [\[Link\]](#)
- **Z. He**, N. Mokhberian, K. Lerman. “Infusing Knowledge from Wikipedia to Enhance Stance Detection.” *WASSA @ EMNLP*, 2022. [\[Link\]](#)
- **Z. He**, L. Tavabi, K. Lerman, M. Soleymani. “Speaker Turn Modeling for Dialogue Act Classification.” *Findings of EMNLP*, 2021. [\[Link\]](#)
- **Z. He**, W. Hao, W.T. Lu, C. Chen, K. Lerman, X. Song. “ALCAP: Alignment-Augmented Music Captioner.” *EMNLP*, 2023. [\[Link\]](#)
- H. Nguyen, **Z. He**, et al. “Smoothing Out Hallucinations: Mitigating LLM Hallucination with Smoothed Knowledge Distillation.” *arXiv*, 2025. [\[Link\]](#)
- **Z. He**, J. May, K. Lerman. “CPL-NoViD: Context-aware Prompt-based Learning for Norm Violation Detection in Online Communities.” *ICWSM*, 2024. [\[Link\]](#)

HONORS

- Best Paper Runner-up, SeT LLM @ ICLR 2024
- Annenberg Fellowship, University of Southern California, 2019
- National Scholarship (top 1.2%), Beijing University of Posts and Telecommunications, 2018

PROFESSIONAL SERVICE

Reviewer: ACL Rolling Review (11 cycles, 2023–2026), COLING 2024/2025, ICWSM 2024, TrustNLP @ NAACL/ACL 2024–2026, R2HCAI @ AAI 2023. Program Committee: CCS 2021.

SKILLS

Languages	Python, C++, CUDA
ML / RL	PyTorch, Transformers, DeepSpeed, FSDP, vLLM, TRL
Data	NumPy, Pandas, Scikit-Learn, Weights & Biases